

QUESTION: 1

You have an application named App1 that reads the data in an Azure Cosmos DB Core (SQL) API account. App1 runs the same read queries every minute. The default consistency level for the account is set to eventual. You discover that every query consumes request units (RUs) instead of using the cache. You verify the IntegratedCacheItemHitRate metric and the IntegratedCacheQueryHitRate metric. Both metrics have values of 0. You verify that the dedicated gateway cluster is provisioned and used in the connection string. You need to ensure that App1 uses the Azure Cosmos DB integrated cache.

What should you configure?

- A. the indexing policy of the Azure Cosmos DB container
- B. the consistency level of the requests from App1
- C. the connectivity mode of the App1 CosmosClient
- D. the default consistency level of the Azure Cosmos DB account

QUESTION: 1

You have an application named App1 that reads the data in an Azure Cosmos DB Core (SQL) API account. App1 runs the same read queries every minute. The default consistency level for the account is set to eventual. You discover that every query consumes request units (RUs) instead of using the cache. You verify the IntegratedCacheItemHitRate metric and the IntegratedCacheQueryHitRate metric. Both metrics have values of 0. You verify that the dedicated gateway cluster is provisioned and used in the connection string. You need to ensure that App1 uses the Azure Cosmos DB integrated cache.

What should you configure?

- A. the indexing policy of the Azure Cosmos DB container
- B. the consistency level of the requests from App1
- C. the connectivity mode of the App1 CosmosClient
- D. the default consistency level of the Azure Cosmos DB account

QUESTION: 2

HOTSPOT

You provision Azure resources by using the following Azure Resource Manager (ARM) template.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
The alert will be triggered when an Azure Cosmos DB key is used	<input type="radio"/>	<input type="radio"/>
Two alert actions will be performed when the alert is triggered	<input type="radio"/>	<input type="radio"/>
The alert will be triggered when an item that has a new partition key value is created	<input type="radio"/>	<input type="radio"/>

QUESTION: 2

HOTSPOT

You provision Azure resources by using the following Azure Resource Manager (ARM) template.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
The alert will be triggered when an Azure Cosmos DB key is used	<input type="radio"/>	<input checked="" type="radio"/>
Two alert actions will be performed when the alert is triggered	<input type="radio"/>	<input checked="" type="radio"/>
The alert will be triggered when an item that has a new partition key value is created	<input checked="" type="radio"/>	<input type="radio"/>

QUESTION: 3**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) API account named account1 that has the disableKeyBasedMetadataWriteAccess property enabled. You are developing an app named App1 that will be used by a user named DevUser1 to create containers in account1. DevUser1 has a non-privileged user account in the Azure Active Directory (Azure AD) tenant. You need to ensure that DevUser1 can use App1 to create containers in account1.

What should you do? (To answer, select the appropriate options in the answer area.)

Answer Area

Grant permissions to create containers by using:

▼
Account keys
Resource tokens
Role-based access control (RBAC)

Create containers by using the:

▼
Azure AD Graph API
Azure Resource Manager API
SQL (Core) API

QUESTION: 3**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) API account named account1 that has the disableKeyBasedMetadataWriteAccess property enabled. You are developing an app named App1 that will be used by a user named DevUser1 to create containers in account1. DevUser1 has a non-privileged user account in the Azure Active Directory (Azure AD) tenant. You need to ensure that DevUser1 can use App1 to create containers in account1.

What should you do? (To answer, select the appropriate options in the answer area.)

Answer Area

Grant permissions to create containers by using:

Account keys
Resource tokens
Role-based access control (RBAC)

Create containers by using the:

Azure AD Graph API
Azure Resource Manager API
SQL (Core) API

QUESTION: 4**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) account that has a single write region in West Europe. You run the following Azure CLI script.

```
az cosmosdb update -n $accountName -g $resourceGroupName \
--locations regionName='West Europe' failoverPriority=0 isZoneRedundant=False \
--locations regionName='North Europe' failoverPriority=1 isZoneRedundant=False

az cosmosdb failover-priority-change -n $accountName -g $resourceGroupName \
--failover-policies 'North Europe=0' 'West Europe=1'
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
After running the script, there will be an instance of Azure Cosmos DB in North Europe that is writable	<input type="radio"/>	<input type="radio"/>
After running the script, the Azure Cosmos DB instance in West Europe will be writable	<input type="radio"/>	<input type="radio"/>
The cost of the Azure Cosmos DB account is unaffected by running the script	<input type="radio"/>	<input type="radio"/>

QUESTION: 4**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) account that has a single write region in West Europe. You run the following Azure CLI script.

```
az cosmosdb update -n $accountName -g $resourceGroupName \
--locations regionName='West Europe' failoverPriority=0 isZoneRedundant=False \
--locations regionName='North Europe' failoverPriority=1 isZoneRedundant=False

az cosmosdb failover-priority-change -n $accountName -g $resourceGroupName \
--failover-policies 'North Europe=0' 'West Europe=1'
```

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
After running the script, there will be an instance of Azure Cosmos DB in North Europe that is writable	<input checked="" type="radio"/>	<input type="radio"/>
After running the script, the Azure Cosmos DB instance in West Europe will be writable	<input type="radio"/>	<input checked="" type="radio"/>
The cost of the Azure Cosmos DB account is unaffected by running the script	<input checked="" type="radio"/>	<input type="radio"/>

QUESTION: 5

You are developing an application that will use an Azure Cosmos DB Core (SQL) API account as a data source. You need to create a report that displays the top five most ordered fruits as shown in the following table.

Name	Type	Orders
apple	fruit	1,000
orange	fruit	600
banana	fruit, exotic	400
plum	fruit.	300
mango	fruit, exotic	200

A collection that contains aggregated data already exists. The following is a sample document:

```
{  
  "name": "apple",  
  "type": ["fruit", "exotic"],  
  "orders": 10000  
}
```

Which two queries can you use to retrieve data for the report? (Select two.)

A.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.orders,i.types
```

B.

```
SELECT TOP i.name, i.types, i.orders.  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.orders DESC
```

C.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.types DESC
```

D.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE ARRAY_CONTAINS(i.types, {name: 'fruit'})  
ORDER BY i.orders DESC
```

- A. Option A
- B. Option B
- C. Option C
- D. Option D

A.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.orders,i.types
```

B.

```
SELECT TOP i.name, i.types, i.orders.  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.orders DESC
```

C.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE EXISTS(SELECT VALUE t FROM t IN i.types WHERE t.name = 'fruit')  
ORDER BY i.types DESC
```

D.

```
SELECT TOP i.name, i.types, i.orders  
FROM items i  
WHERE ARRAY_CONTAINS(i.types, {name: 'fruit'})  
ORDER BY i.orders DESC
```

A. Option A

B. Option B

C. Option C

D. Option D

QUESTION: 6**HOTSPOT**

You have a database in an Azure Cosmos DB Core (SQL) API account. You plan to create a container that will store employee data for 5,000 small businesses. Each business will have up to 25 employees. Each employee item will have an emailAddress value. You need to ensure that the emailAddress value for each employee within the same company is unique.

To what should you set the partition key and the unique key? (To answer, select the appropriate options in the answer area.)

Answer Area

Partition key

▼
companyId
companyId+emailAddress
emailAddress
employeeId

Unique key

▼
companyId
emailAddress
employeeId

QUESTION: 6**HOTSPOT**

You have a database in an Azure Cosmos DB Core (SQL) API account. You plan to create a container that will store employee data for 5,000 small businesses. Each business will have up to 25 employees. Each employee item will have an emailAddress value. You need to ensure that the emailAddress value for each employee within the same company is unique.

To what should you set the partition key and the unique key? (To answer, select the appropriate options in the answer area.)

Answer Area

Partition key

▼
companyId
companyId+emailAddress
emailAddress
employeeId

Unique key

▼
companyId
emailAddress
employeeId

QUESTION: 7**HOTSPOT**

You have a container named container1 in an Azure Cosmos DB Core (SQL) API account. The container1 container has 120 GB of data. The following is a sample of a document in container1.

```
{  
    "customerId" : "5425",  
    "orderId" : "9d7816e6-f401-42ba-ad05-0e03de35c0b8",  
    "orderDate" : "2019-05-03",  
    "orderDetails" : []  
}
```

The orderId property is used as the partition key.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.orderDate = "2019-05-03"</pre>	<input type="radio"/>	<input type="radio"/>
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.customerId = "5425"</pre>	<input type="radio"/>	<input type="radio"/>
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.orderDate = "2019-05-03" and c.orderId = "9d7816e6-f401-42ba-ad05-0e03de35c0b8"</pre>	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.orderDate = "2019-05-03"</pre>	<input checked="" type="radio"/>	<input type="radio"/>
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.customerId = "5425"</pre>	<input checked="" type="radio"/>	<input type="radio"/>
If you run the following query, the query will run as a cross-partition query <pre>SELECT * FROM c where c.orderDate = "2019-05-03" and c.orderId = "9d7816e6-f401-42ba-ad05-0e03de35c0b8"</pre>	<input type="radio"/>	<input checked="" type="radio"/>

QUESTION: 8

You are designing an Azure Cosmos DB Core (SQL) API solution to store data from IoT devices. Writes from the devices will occur every second. The following is a sample of the data.

```
{  
  "id" : "03c1ca5a-db18-4231-908f-09a9bc7a7c3e",  
  "deviceManufacturer" : "Contoso, Ltd",  
  "deviceId" : "f460df85-799f-4d58-b051-67561b4993c6",  
  "timestamp" : "2021-09-19T13:47:45",  
  "sensor1Value" : true,  
  "sensor2Value" : "75",  
  "sensor3Value" : "4554",  
  "sensor4Value" : "454",  
  "sensor5Value" : "42128"  
}
```

You need to select a partition key that meets the following requirements for writes:

- Minimizes the partition skew
- Avoids capacity limits
- Avoids hot partitions

What should you do?

- A. Use timestamp as the partition key.
- B. Create a new synthetic key that contains deviceId and sensor1Value.
- C. Create a new synthetic key that contains deviceId and deviceManufacturer.
- D. Create a new synthetic key that contains deviceId and a random number.

QUESTION: 8

You are designing an Azure Cosmos DB Core (SQL) API solution to store data from IoT devices. Writes from the devices will occur every second. The following is a sample of the data.

```
{  
  "id" : "03c1ca5a-db18-4231-908f-09a9bc7a7c3e",  
  "deviceManufacturer" : "Contoso, Ltd",  
  "deviceId" : "f460df85-799f-4d58-b051-67561b4993c6",  
  "timestamp" : "2021-09-19T13:47:45",  
  "sensor1Value" : true,  
  "sensor2Value" : "75",  
  "sensor3Value" : "4554",  
  "sensor4Value" : "454",  
  "sensor5Value" : "42128"  
}
```

You need to select a partition key that meets the following requirements for writes:

- Minimizes the partition skew
- Avoids capacity limits
- Avoids hot partitions

What should you do?

- A. Use timestamp as the partition key.
- B. Create a new synthetic key that contains deviceId and sensor1Value.
- C. Create a new synthetic key that contains deviceId and deviceManufacturer.
- D. Create a new synthetic key that contains deviceId and a random number.

QUESTION: 9

You maintain a relational database for a book publisher. The database contains the following tables.

The most common query lists the books for a given authorId. You need to develop a non-relational data model for Azure Cosmos DB Core (SQL) API that will replace the relational database. The solution must minimize latency and read operation costs.

What should you include in the solution?

- A. Create a container for Author and a container for Book. In each Author document, embed bookId for each book by the author. In each Book document embed authorId of each author.
- B. Create Author, Book, and Bookauthorlnk documents in the same container.
- C. Create a container that contains a document for each Author and a document for each Book.
- D. Create a container for Author and a container for Book. In each Author document and Book document embed the data from Bookauthorlnk.

Name	Column
Author	authorId (primary key)
	fullname
	address
	contactinfo
Book	bookId (primary key)
	isbn
	title
	genre
Bookauthorlnk	authorId (foreign key)
	bookId (foreign key)

QUESTION: 9

You maintain a relational database for a book publisher. The database contains the following tables.

The most common query lists the books for a given authorId. You need to develop a non-relational data model for Azure Cosmos DB Core (SQL) API that will replace the relational database. The solution must minimize latency and read operation costs.

What should you include in the solution?

- A. Create a container for Author and a container for Book. In each Author document, embed bookId for each book by the author. In each Book document embed authorId of each author.
- B. Create Author, Book, and Bookauthorlnk documents in the same container.
- C. Create a container that contains a document for each Author and a document for each Book.
- D. Create a container for Author and a container for Book. In each Author document and Book document embed the data from Bookauthorlnk.

Name	Column
Author	authorId (primary key)
	fullname
	address
	contactinfo
Book	bookId (primary key)
	isbn
	title
	genre
Bookauthorlnk	authorId (foreign key)
	bookId (foreign key)

QUESTION: 10

You have an Azure Cosmos DB Core (SQL) API account. You run the following query against a container in the account.

```
SELECT  
IS_NUMBER("1234") AS A,  
IS_NUMBER(1234) AS B,  
IS_NUMBER({prop: 1234}) AS C
```

What is the output of the query?

- A. [{"A": false, "B": true, "C": false}]
- B. [{"A": true, "B": false, "C": true}]
- C. [{"A": true, "B": true, "C": false}]
- D. [{"A": true, "B": true, "C": true}]

QUESTION: 10

You have an Azure Cosmos DB Core (SQL) API account. You run the following query against a container in the account.

```
SELECT  
IS_NUMBER("1234") AS A,  
IS_NUMBER(1234) AS B,  
IS_NUMBER({prop: 1234}) AS C
```

What is the output of the query?

- A. [{"A": false, "B": true, "C": false}]
- B. [{"A": true, "B": false, "C": true}]
- C. [{"A": true, "B": true, "C": false}]
- D. [{"A": true, "B": true, "C": true}]

QUESTION: 11

You need to implement a trigger in Azure Cosmos DB Core (SQL) API that will run before an item is inserted into a container.

Which two actions should you perform to ensure that the trigger runs? (Select two.)

- A. Append pre to the name of the JavaScript function trigger.
- B. For each create request, set the access condition in RequestOptions.
- C. Register the trigger as a pre-trigger.
- D. For each create request, set the consistency level to session in RequestOptions.
- E. For each create request, set the trigger name in RequestOptions.

QUESTION: 11

You need to implement a trigger in Azure Cosmos DB Core (SQL) API that will run before an item is inserted into a container.

Which two actions should you perform to ensure that the trigger runs? (Select two.)

- A. Append pre to the name of the JavaScript function trigger.
- B. For each create request, set the access condition in RequestOptions.
- C. Register the trigger as a pre-trigger.
- D. For each create request, set the consistency level to session in RequestOptions.
- E. For each create request, set the trigger name in RequestOptions.

QUESTION: 12

HOTSPOT

You have a container in an Azure Cosmos DB Core (SQL) API account. You need to use the Azure Cosmos DB SDK to replace a document by using optimistic concurrency.

What should you include in the code? (To answer, select the appropriate options in the answer area.)

Answer Area

RequestOptions property to set:

	▼
AccessCondition	
ConsistencyLevel	
SessionToken	

Document property that will be compared:

	▼
_etag	
_id	
_rid	

QUESTION: 12

HOTSPOT

You have a container in an Azure Cosmos DB Core (SQL) API account. You need to use the Azure Cosmos DB SDK to replace a document by using optimistic concurrency.

What should you include in the code? (To answer, select the appropriate options in the answer area.)

Answer Area

RequestOptions property to set:

	▼
AccessCondition	
ConsistencyLevel	
SessionToken	

Document property that will be compared:

	▼
_etag	
_id	
_rid	

QUESTION: 13**HOTSPOT**

You are creating a database in an Azure Cosmos DB Core (SQL) API account. The database will be used by an application that will provide users with the ability to share online posts. Users will also be able to submit comments on other users' posts. You need to store the data shown in the following table.

Type	Description
Users	Information about a user who will use the application
Posts	Text of up to 1,000 characters that a user will share with other users
Comments	Text of up to 280 characters that users will submit as a comment on a post
Interests	Information about a user's interests

The application has the following characteristics:

- Users can submit an unlimited number of posts.
- The average number of posts submitted by a user will be more than 1,000.
- Posts can have an unlimited number of comments from different users.
- The average number of comments per post will be 100, but many posts will exceed 1,000 comments.
- Users will be limited to having a maximum of 20 interests.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
If you embed the posts data into the users data instead of creating a separate document for each post, you will increase the write operation costs for new posts	<input type="radio"/>	<input type="radio"/>
If you embed the comments data into the posts data instead of creating a separate document for each comment you will increase the write operation costs for new comments	<input type="radio"/>	<input type="radio"/>
If you embed the interests data into the users data instead of creating a separate document for each interest, you will increase the read operation costs for displaying the users and their associated interests	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
If you embed the posts data into the users data instead of creating a separate document for each post, you will increase the write operation costs for new posts	<input checked="" type="radio"/>	<input type="radio"/>
If you embed the comments data into the posts data instead of creating a separate document for each comment you will increase the write operation costs for new comments	<input checked="" type="radio"/>	<input type="radio"/>
If you embed the interests data into the users data instead of creating a separate document for each interest, you will increase the read operation costs for displaying the users and their associated interests	<input type="radio"/>	<input checked="" type="radio"/>

QUESTION: 14**DRAG DROP**

You have an app that stores data in an Azure Cosmos DB Core (SQL) API account. The app performs queries that return large result sets. You need to return a complete result set to the app by using pagination. Each page of results must return 80 items.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Actions

Configure MaxItemCount in QueryRequestOptions

Run the query and provide a continuation token

Configure MaxBufferedItemCount in
QueryRequestOptions

Append the results to a variable

Run the query and increment MaxItemCount

Answer Area

QUESTION: 14

DRAG DROP

You have an app that stores data in an Azure Cosmos DB Core (SQL) API account. The app performs queries that return large result sets. You need to return a complete result set to the app by using pagination. Each page of results must return 80 items.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Actions

Configure MaxItemCount in QueryRequestOptions

Run the query and provide a continuation token

Configure MaxBufferedItemCount in
QueryRequestOptions

Append the results to a variable

Run the query and increment MaxItemCount

Answer Area

Configure MaxItemCount in QueryRequestOptions

Run the query and provide a continuation token

Append the results to a variable



Yes or No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure an Azure Monitor alert to trigger the function.

Does this meet the goal?

- A. Yes
- B. No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure an Azure Monitor alert to trigger the function.

Does this meet the goal?

A. Yes

B. No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure the function to have an Azure CosmosDB trigger.

Does this meet the goal?

- A. Yes
- B. No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure the function to have an Azure CosmosDB trigger.

Does this meet the goal?

- A. Yes
- B. No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure an application to use the change feed processor to read the change feed and you configure the application to trigger the function.

Does this meet the goal?

- A. Yes
- B. No

QUESTION: 15

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You have an Azure Cosmos DB Core (SQL) API account named account 1 that uses autoscale throughput. You need to run an Azure function when the normalized request units per second for a container in account1 exceeds a specific value.

Solution: You configure an application to use the change feed processor to read the change feed and you configure the application to trigger the function.

Does this meet the goal?

- A. Yes
- B. No

Yes or No

QUESTION: 18**DRAG DROP**

You have an Azure Cosmos DB Core (SQL) API account that is configured for multi-region writes. The account contains a database that has two containers named container1 and container2. The following is a sample of a document in container1:

```
{  
  "customerId": 1234,  
  "firstName": "John",  
  "lastName": "Smith",  
  "policyYear": 2021  
}
```

The following is a sample of a document in container2:

```
{  
  "gpsId": 1234,  
  "latitude": 38.8951,  
  "longitude": -77.0364  
}
```

You need to configure conflict resolution to meet the following requirements:

- For container1 you must resolve conflicts by using the highest value for policyYear.
- For container2 you must resolve conflicts by accepting the distance closest to latitude: 40.730610 and longitude: -73.935242.
- Administrative effort must be minimized to implement the solution.

What should you configure for each container? (To answer, drag the appropriate configurations to the correct containers. Each configuration may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Configurations

Last Write Wins (default) mode

Merge Procedures (custom) mode

An application that reads from the conflicts feed

Answer Area

Container1:



Container2:



You need to configure conflict resolution to meet the following requirements:

- For container1 you must resolve conflicts by using the highest value for policyYear.
- For container2 you must resolve conflicts by accepting the distance closest to latitude: 40.730610 and longitude: -73.935242.
- Administrative effort must be minimized to implement the solution.

What should you configure for each container? (To answer, drag the appropriate configurations to the correct containers. Each configuration may be used once, more than once, or not at all.)

You may need to drag the split bar between panes or scroll to view content.)

Configurations

Last Write Wins (default) mode

Merge Procedures (custom) mode

An application that reads from the conflicts feed

Answer Area

Container1:

Last Write Wins (default) mode

Container2:

Merge Procedures (custom) mode

QUESTION: 19

You have a container in an Azure Cosmos DB Core (SQL) API account. The container stores telemetry data from IoT devices. The container uses telemetryId as the partition key and has a throughput of 1,000 request units per second (RU/s). Approximately 5,000 IoT devices submit data every five minutes by using the same telemetryId value. You have an application that performs analytics on the data and frequently reads telemetry data for a single IoT device to perform trend analysis. The following is a sample of a document in the container.

```
{  
    "id" : "9ccf1906-2a30-4dc0-9644-2185f5dcbbd7",  
    "deviceId" : "bba6fe24-6d97-4935-8d58-36baa4b8a0e1",  
    "telemetryId" : "9d7816e6-f401-42ba-ad05-0e03de35c0b8",  
    "date" : "2019-05-03",  
    "time" : "13:05",  
    "temp" : "21"  
}
```

You need to reduce the amount of request units (RUs) consumed by the analytics application.

What should you do?

- A. Decrease the offerThroughput value for the container.
- B. Increase the offerThroughput value for the container.
- C. Move the data to a new container that has a partition key of deviceId.
- D. Move the data to a new container that uses a partition key of date.

QUESTION: 19

You have a container in an Azure Cosmos DB Core (SQL) API account. The container stores telemetry data from IoT devices. The container uses telemetryId as the partition key and has a throughput of 1,000 request units per second (RU/s). Approximately 5,000 IoT devices submit data every five minutes by using the same telemetryId value. You have an application that performs analytics on the data and frequently reads telemetry data for a single IoT device to perform trend analysis. The following is a sample of a document in the container.

```
{  
    "id" : "9ccf1906-2a30-4dc0-9644-2185f5dcbbd7",  
    "deviceId" : "bba6fe24-6d97-4935-8d58-36baa4b8a0e1",  
    "telemetryId" : "9d7816e6-f401-42ba-ad05-0e03de35c0b8",  
    "date" : "2019-05-03",  
    "time" : "13:05",  
    "temp" : "21"  
}
```

You need to reduce the amount of request units (RUs) consumed by the analytics application.

What should you do?

- A. Decrease the offerThroughput value for the container.
- B. Increase the offerThroughput value for the container.
- C. Move the data to a new container that has a partition key of deviceId.
- D. Move the data to a new container that uses a partition key of date.

QUESTION: 20**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) API account named storage1 that uses provisioned throughput capacity mode. The storage1 account contains the databases shown in the following table.

Name	Throughput	Max request units per second (RU/s)	Geo-redundancy	Multi-region writes	Number of regions
db1	Autoscale	5,000	Disabled	Disabled	1
db2	Autoscale	8,000	Enabled	Enabled	3

The databases contain the containers shown in the following table.

Name	Database	Throughput
cn01	db1	Container - autoscale maximum RU/s of 10,000
cn02	db1	Database
cn03	db1	Database
cn04	db1	Database
cn05	db1	Database
cn11	db2	Database
cn12	db2	Database
cn13	db2	Database
cn14	db2	Database
cn15	db2	Database
cn16	db2	Database
cn17	db2	Database
cn18	db2	Database

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
At a minimum, you will be billed for 4,000 RU/s per hour for db1	<input type="radio"/>	<input type="radio"/>
The maximum throughput that can be consumed by cn11 is 400 RU/s	<input type="radio"/>	<input type="radio"/>
To db2, you can add a new container that uses database throughput	<input type="radio"/>	<input type="radio"/>

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
At a minimum, you will be billed for 4,000 RU/s per hour for db1	<input type="radio"/>	<input checked="" type="radio"/>
The maximum throughput that can be consumed by cn11 is 400 RU/s	<input type="radio"/>	<input checked="" type="radio"/>
To db2, you can add a new container that uses database throughput	<input checked="" type="radio"/>	<input type="radio"/>

QUESTION: 21

HOTSPOT

You have a database named telemetry in an Azure Cosmos DB Core (SQL) API account that stores IoT data. The database contains two containers named readings and devices.

Documents in readings have the following structure.

- id
- deviceid
- timestamp
- ownerid
- measures (array)
- type
- value
- metricid

Documents in devices have the following structure.

- id
- deviceid
- owner
- ownerid
- emailaddress
- name
- brand
- model

Answer Area

Statements	Yes	No
To return for all devices owned by a specific emailaddress, multiple queries must be performed	<input type="radio"/>	<input type="radio"/>
To return deviceid, ownerid, timestamp, and value for a specific metricid, a join must be performed	<input type="radio"/>	<input type="radio"/>
To return deviceid, ownerid, emailaddress, and model, a join must be performed	<input type="radio"/>	<input type="radio"/>

Answer Area

Statements	Yes	No
To return for all devices owned by a specific emailaddress, multiple queries must be performed	<input checked="" type="radio"/>	<input type="radio"/>
To return deviceid, ownerid, timestamp, and value for a specific metricid, a join must be performed	<input type="radio"/>	<input checked="" type="radio"/>
To return deviceid, ownerid, emailaddress, and model, a join must be performed	<input type="radio"/>	<input checked="" type="radio"/>

QUESTION: 22

The settings for a container in an Azure Cosmos DB Core (SQL) API account are configured as shown in the following exhibit.

Settings

Indexing Policy

Time to Live

- Off
- On (no default)
- On

Which statement describes the configuration of the container?

- A. All items will be deleted after one year.
- B. Items stored in the collection will be retained always, regardless of the items time to live value.
- C. Items stored in the collection will expire only if the item has a time to live value.
- D. All items will be deleted after one hour.

Geospatial Configuration

- Geography
- Geometry

Partition key

```
/productName
```

QUESTION: 22

The settings for a container in an Azure Cosmos DB Core (SQL) API account are configured as shown in the following exhibit.

Settings

Indexing Policy

Time to Live

- Off
- On (no default)
- On

Which statement describes the configuration of the container?

- A. All items will be deleted after one year.
- B. Items stored in the collection will be retained always, regardless of the items time to live value.
- C. Items stored in the collection will expire only if the item has a time to live value.
- D. All items will be deleted after one hour.

Geospatial Configuration

- Geography
- Geometry

Partition key

```
/productName
```

QUESTION: 23

You have an Azure Cosmos DB Core (SQL) API account that uses a custom conflict resolution policy. The account has a registered merge procedure that throws a runtime exception. The runtime exception prevents conflicts from being resolved. You need to use an Azure function to resolve the conflicts.

What should you use?

- A. a function that pulls items from the conflicts feed and is triggered by a timer trigger
- B. a function that receives items pushed from the change feed and is triggered by an Azure Cosmos DB trigger
- C. a function that pulls items from the change feed and is triggered by a timer trigger
- D. a function that receives items pushed from the conflicts feed and is triggered by an Azure Cosmos DB trigger

QUESTION: 23

You have an Azure Cosmos DB Core (SQL) API account that uses a custom conflict resolution policy. The account has a registered merge procedure that throws a runtime exception. The runtime exception prevents conflicts from being resolved. You need to use an Azure function to resolve the conflicts.

What should you use?

- A. a function that pulls items from the conflicts feed and is triggered by a timer trigger
- B. a function that receives items pushed from the change feed and is triggered by an Azure Cosmos DB trigger
- C. a function that pulls items from the change feed and is triggered by a timer trigger
- D. a function that receives items pushed from the conflicts feed and is triggered by an Azure Cosmos DB trigger

QUESTION: 24

The following is a sample of a document in orders.

```
{  
    "orderId" : "d4a91979b-5ead-43a3-b851-add9a71ac4b6",  
    "customerId" : "f6e39103-bdc7-4346-9cfb-45daa4b2becf",  
    "orderDate" : "2021-09-29",  
    "orderItems" : [  
        {  
            "itemId" : "6c30412f-3cd7-4cab-813c-05942345720d",  
            "name" : "blue pen",  
            "type" : "pens",  
            "count" : 10,  
        },  
        ...  
    ],  
    "total" : 12345,  
    "status" : "ordered"  
}
```

The orders container uses customerId as the partition key. You need to provide a report of the total items ordered per month by item type. The solution must meet the following requirements:

- Ensure that the report can run as quickly as possible.
- Minimize the consumption of request units (RUs).

What should you do?

- A. Configure the report to query orders by using a SQL query.
- B. Configure the report to query a new aggregate container. Populate the aggregates by using the change feed.
- C. Configure the report to query orders by using a SQL query through a dedicated gateway.
- D. Configure the report to query a new aggregate container. Populate the aggregates by using SQL queries that run daily.

The orders container uses customerId as the partition key. You need to provide a report of the total items ordered per month by item type. The solution must meet the following requirements:

- Ensure that the report can run as quickly as possible.
- Minimize the consumption of request units (RUs).

What should you do?

- A. Configure the report to query orders by using a SQL query.
- B. Configure the report to query a new aggregate container. Populate the aggregates by using the change feed.
- C. Configure the report to query orders by using a SQL query through a dedicated gateway.
- D. Configure the report to query a new aggregate container. Populate the aggregates by using SQL queries that run daily.

QUESTION: 25

You configure multi-region writes for account1. You need to ensure that App1 supports the new configuration for account1. The solution must meet the business requirements and the product catalog requirements.

What should you do?

- A. Create a private endpoint connection.
- B. Increase the number of request units per second (RU/s) allocated to the con-product and conproductVendor containers.
- C. Modify the connection policy of App1.
- D. Set the default consistency level of account1 to bounded staleness.

QUESTION: 25

You configure multi-region writes for account1. You need to ensure that App1 supports the new configuration for account1. The solution must meet the business requirements and the product catalog requirements.

What should you do?

- A. Create a private endpoint connection.
- B. Increase the number of request units per second (RU/s) allocated to the con-product and conproductVendor containers.
- C. Modify the connection policy of App1.
- D. Set the default consistency level of account1 to bounded staleness.

QUESTION: 26

You are troubleshooting the current issues caused by the application updates.

Which action can address the application updates issue without affecting the functionality of the application?

- A. Enable time to live for the con-product container.
- B. Set the default consistency level of account1 to strong.
- C. Set the default consistency level of account1 to bounded staleness.
- D. Add a custom indexing policy to the con-product container.

QUESTION: 26

You are troubleshooting the current issues caused by the application updates.

Which action can address the application updates issue without affecting the functionality of the application?

- A. Enable time to live for the con-product container.
- B. Set the default consistency level of account1 to strong.
- C. Set the default consistency level of account1 to bounded staleness.
- D. Add a custom indexing policy to the con-product container.

QUESTION: 27

You need to select the partition key for con-iot1. The solution must meet the IoT telemetry requirements.

What should you select?

- A. the timestamp
- B. the humidity
- C. the temperature
- D. the device ID

QUESTION: 27

You need to select the partition key for con-iot1. The solution must meet the IoT telemetry requirements.

What should you select?

- A. the timestamp
- B. the humidity
- C. the temperature
- D. the device ID

QUESTION: 28

You need to identify which connectivity mode to use when implementing App2. The solution must support the planned changes and meet the business requirements.

Which connectivity mode should you identify?

- A. Direct mode over HTTPS
- B. Gateway mode (using HTTPS)
- C. Direct mode over TCP

QUESTION: 28

You need to identify which connectivity mode to use when implementing App2. The solution must support the planned changes and meet the business requirements.

Which connectivity mode should you identify?

- A. Direct mode over HTTPS
- B. Gateway mode (using HTTPS)
- C. Direct mode over TCP

QUESTION: 29

You need to provide a solution for the Azure Functions notifications following updates to con-product. The solution must meet the business requirements and the product catalog requirements.

Which two actions should you perform? (Select two.)

- A. Configure the trigger for each function to use a different leaseCollectionPrefix
- B. Configure the trigger for each function to use the same leaseCollectionName
- C. Configure the trigger for each function to use a different leaseCollectionName
- D. Configure the trigger for each function to use the same leaseCollectionPrefix

QUESTION: 29

You need to provide a solution for the Azure Functions notifications following updates to con-product. The solution must meet the business requirements and the product catalog requirements.

Which two actions should you perform? (Select two.)

- A. Configure the trigger for each function to use a different leaseCollectionPrefix
- B. Configure the trigger for each function to use the same leaseCollectionName
- C. Configure the trigger for each function to use a different leaseCollectionName
- D. Configure the trigger for each function to use the same leaseCollectionPrefix

QUESTION: 30

You have the following query.

```
SELECT * FROM  
WHERE c.sensor = "TEMP1"  
AND c.value < 22  
AND c.timestamp >= 1619146031231
```

You need to recommend a composite index strategy that will minimize the request units (RUs) consumed by the query.

What should you recommend?

- A. a composite index for (sensor ASC, value ASC) and a composite index for (sensor ASC, timestamp ASC)
- B. a composite index for (sensor ASC, value ASC, timestamp ASC) and a composite index for (sensor DESC, value DESC, timestamp DESC)
- C. a composite index for (value ASC, sensor ASC) and a composite index for (timestamp ASC, sensor ASC)
- D. a composite index for (sensor ASC, value ASC, timestamp ASC)

QUESTION: 30

You have the following query.

```
SELECT * FROM  
WHERE c.sensor = "TEMP1"  
AND c.value < 22  
AND c.timestamp >= 1619146031231
```

You need to recommend a composite index strategy that will minimize the request units (RUs) consumed by the query.

What should you recommend?

- A. a composite index for (sensor ASC, value ASC) and a composite index for (sensor ASC, timestamp ASC)
- B. a composite index for (sensor ASC, value ASC, timestamp ASC) and a composite index for (sensor DESC, value DESC, timestamp DESC)
- C. a composite index for (value ASC, sensor ASC) and a composite index for (timestamp ASC, sensor ASC)
- D. a composite index for (sensor ASC, value ASC, timestamp ASC)

QUESTION: 31

You have an Azure Cosmos DB Core (SQL) API account. You configure the diagnostic settings to send all log information to a Log Analytics workspace. You need to identify when the provisioned request units per second (RU/s) for resources within the account were modified.

You write the following query.

```
AzureDiagnostics  
| where Category == "ControlPlaneRequests"
```

What should you include in the query?

- A. | where OperationName startswith "AccountUpdateStart"
- B. | where OperationName startswith "SqlContainersDelete"
- C. | where OperationName startswith "MongoCollectionsThroughputUpdate"
- D. | where OperationName startswith "SqlContainersThroughputUpdate"

QUESTION: 31

You have an Azure Cosmos DB Core (SQL) API account. You configure the diagnostic settings to send all log information to a Log Analytics workspace. You need to identify when the provisioned request units per second (RU/s) for resources within the account were modified.

You write the following query.

```
AzureDiagnostics  
| where Category == "ControlPlaneRequests"
```

What should you include in the query?

- A. | where OperationName startswith "AccountUpdateStart"
- B. | where OperationName startswith "SqlContainersDelete"
- C. | where OperationName startswith "MongoCollectionsThroughputUpdate"
- D. | where OperationName startswith "SqlContainersThroughputUpdate"

QUESTION: 32

You plan to create an Azure Cosmos DB Core (SQL) API account that will use customer-managed keys stored in Azure Key Vault. You need to configure an access policy in Key Vault to allow Azure Cosmos DB access to the keys.

Which three permissions should you enable in the access policy? (Select three.)

- A. Wrap Key
- B. Get
- C. List
- D. Update
- E. Sign
- F. Verify
- G. Unwrap Key

QUESTION: 32

You plan to create an Azure Cosmos DB Core (SQL) API account that will use customer-managed keys stored in Azure Key Vault. You need to configure an access policy in Key Vault to allow Azure Cosmos DB access to the keys.

Which three permissions should you enable in the access policy? (Select three.)

- A. Wrap Key
- B. Get
- C. List
- D. Update
- E. Sign
- F. Verify
- G. Unwrap Key

QUESTION: 33

You need to configure an Apache Kafka instance to ingest data from an Azure Cosmos DB Core (SQL) API account. The data from a container named telemetry must be added to a Kafka topic named IoT. The solution must store the data in a compact binary format.

Which three configuration items should you include in the solution? (Select three.)

- A. "connector.class": "com.azure.cosmos.kafka.connect.source.CosmosDBSourceConnector"
- B. "key.converter": "org.apache.kafka.connect.json.JsonConverter"
- C. "key.converter": "io.confluent.connect.avro.AvroConverter"
- D. "connect.cosmos.containers.topicmap": "iot#telemetry"
- E. "connect.cosmos.containers.topicmap": "iot"
- F. "connector.class": "com.azure.cosmos.kafka.connect.source.CosmosDBSinkConnector"

QUESTION: 33

You need to configure an Apache Kafka instance to ingest data from an Azure Cosmos DB Core (SQL) API account. The data from a container named telemetry must be added to a Kafka topic named IoT. The solution must store the data in a compact binary format.

Which three configuration items should you include in the solution? (Select three.)

- A. "connector.class": "com.azure.cosmos.kafka.connect.source.CosmosDBSourceConnector"
- B. "key.converter": "org.apache.kafka.connect.json.JsonConverter"
- C. "key.converter": "io.confluent.connect.avro.AvroConverter"
- D. "connect.cosmos.containers.topicmap": "iot#telemetry"
- E. "connect.cosmos.containers.topicmap": "iot"
- F. "connector.class": "com.azure.cosmos.kafka.connect.source.CosmosDBSinkConnector"

QUESTION: 34**HOTSPOT**

You plan to deploy two Azure Cosmos DB Core (SQL) API accounts that will each contain a single database. The accounts will be configured as shown in the following table.

Name	Description
development	<ul style="list-style-type: none">• Supports the development of new application features• Used intermittently as needed during development
shipments	<ul style="list-style-type: none">• Captures over 100,000 updates per second generated at unpredictable times throughout the business day• Used with Azure Synapse Link for analytics

How should you provision the containers within each account to minimize costs? (To answer, select the appropriate options in the answer area.)

development:

- Serverless capacity mode
- Provisioned throughput capacity mode and manual throughput
- Provisioned throughput capacity mode and autoscale throughput

shipments:

- Serverless capacity mode
- Provisioned throughput capacity mode and manual throughput
- Provisioned throughput capacity mode and autoscale throughput

How should you provision the containers within each account to minimize costs? (To answer, select the appropriate options in the answer area.)

development:

- Serverless capacity mode
- Provisioned throughput capacity mode and manual throughput
- Provisioned throughput capacity mode and autoscale throughput

shipments:

- Serverless capacity mode
- Provisioned throughput capacity mode and manual throughput
- Provisioned throughput capacity mode and autoscale throughput

QUESTION: 35**HOTSPOT**

You have an Azure Cosmos DB Core (SQL) API account named account1. You have the Azure virtual networks and subnets shown in the following table.

Subnet	Network	IP address range	Virtual machine
subnet1	vnet1	10.0.0.0/24	VM1
subnet2	vnet1	10.0.1.0/24	VM2
subnet3	vnet2	10.1.0.0/24	VM3

The vnet1 and vnet2 networks are connected by using a virtual network peer. The Firewall and virtual network settings for account1 are configured as shown in the exhibit.

Allow access from

- All networks Selected networks

Configure network security for your Azure Cosmos DB account. [Learn more.](#)

Virtual networks

Secure your Azure Cosmos DB account with virtual networks. [+ Add existing virtual network](#) [+Add new virtual network](#)

Virtual Network	Subnet	Address range	Endpoint Status
vnet1	1	10.0.0.0/16	
	vnet1.subnet1	10.0.1.0/24	✓ Enabled

Firewall

Add IP ranges to allow access from the internet or your on-premises networks. [+Add my current IP](#)

ⓘ

IP(Single IPv4 or CIDR range)

Exceptions

- Accept connections from within public Azure datacenters ⓘ
- Allow access from Azure Portal ⓘ

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Statements	Yes	No
VM1 can access account 1	<input type="radio"/>	<input type="radio"/>
VM2 can access account 1	<input type="radio"/>	<input type="radio"/>
VM3 can access account 1	<input type="radio"/>	<input type="radio"/>

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Statements	Yes	No
VM1 can access account 1	<input checked="" type="radio"/>	<input type="radio"/>
VM2 can access account 1	<input type="radio"/>	<input checked="" type="radio"/>
VM3 can access account 1	<input type="radio"/>	<input checked="" type="radio"/>

QUESTION: 36

You have a container named container1 in an Azure Cosmos DB Core (SQL) API account. You need to provide a user named User1 with the ability to insert items into container1 by using role-based access control (RBAC). The solution must use the principle of least privilege.

Which roles should you assign to User1?

- A. CosmosDB Operator only
- B. DocumentDB Account Contributor and Cosmos DB Built-in Data Contributor
- C. DocumentDB Account Contributor only
- D. Cosmos DB Built-in Data Contributor only

QUESTION: 36

You have a container named container1 in an Azure Cosmos DB Core (SQL) API account. You need to provide a user named User1 with the ability to insert items into container1 by using role-based access control (RBAC). The solution must use the principle of least privilege.

Which roles should you assign to User1?

- A. CosmosDB Operator only
- B. DocumentDB Account Contributor and Cosmos DB Built-in Data Contributor
- C. DocumentDB Account Contributor only
- D. Cosmos DB Built-in Data Contributor only

QUESTION: 37

You have an existing Azure storage account that stores large volumes of data across multiple containers. You need to copy all data from the existing storage account to a new storage account. The copy process must meet the following requirements:

- Automate data movement.
- Minimize user input required to perform the operation.
- Ensure that the data movement process is recoverable.

What should you use?

- A. AzCopy
- B. Azure Storage Explorer
- C. Azure portal
- D. .NET Storage Client Library

QUESTION: 37

You have an existing Azure storage account that stores large volumes of data across multiple containers. You need to copy all data from the existing storage account to a new storage account. The copy process must meet the following requirements:

- Automate data movement.
- Minimize user input required to perform the operation.
- Ensure that the data movement process is recoverable.

What should you use?

- A. AzCopy
- B. Azure Storage Explorer
- C. Azure portal
- D. .NET Storage Client Library

QUESTION: 38

Which category of SQL statements is used to add, remove, and modify database structures?

- A. Data access language (DAL)
- B. Data manipulation language (DML)
- C. Data control language (DCL)
- D. Data definition language (DDL)

QUESTION: 38

Which category of SQL statements is used to add, remove, and modify database structures?

- A. Data access language (DAL)
- B. Data manipulation language (DML)
- C. Data control language (DCL)
- D. Data definition language (DDL)

QUESTION: 39

You execute the following statement:

```
SELECT EmployeeID, FirstName, DepartmentName  
FROM Employee, Department
```

This type of operation is called a/an:

- A. Intersection
- B. Outer join
- C. Equi-join
- D. Cartesian product

QUESTION: 39

You execute the following statement:

```
SELECT EmployeeID, FirstName, DepartmentName  
FROM Employee, Department
```

This type of operation is called a/an:

- A. Intersection
- B. Outer join
- C. Equi-join
- D. Cartesian product

QUESTION: 40

You have a table named Product that contains the following data.

ProductID	ProductName	CategoryID
3296	Spoon	2222
1114	Chair	4444

The ProductID column is the primary key. The CategoryID column is a foreign key to a separate table named Category. You execute the following statement:

```
INSERT INTO Product  
VALUES (3296, 'Table', 4444)
```

What is the result?

- A. a foreign key constraint violation
- B. a syntax error
- C. a new row in the Product table
- D. a primary key constraint violation
- E. a new row in the Category table

QUESTION: 40

You have a table named Product that contains the following data.

ProductID	ProductName	CategoryID
3296	Spoon	2222
1114	Chair	4444

The ProductID column is the primary key. The CategoryID column is a foreign key to a separate table named Category. You execute the following statement:

```
INSERT INTO Product  
VALUES (3296, 'Table', 4444)
```

What is the result?

- A. a foreign key constraint violation
- B. a syntax error
- C. a new row in the Product table
- D. a primary key constraint violation
- E. a new row in the Category table

QUESTION: 41

Your company has a web app named WebApp1. You use the WebJobs SDK to design a triggered App Service background task that automatically invokes a function in the code every time new data is received in a queue. You are preparing to configure the service processes a queue data item.

Which of the following is the service you should use?

- A. Logic Apps
- B. WebJobs
- C. Flow
- D. Functions

QUESTION: 41

Your company has a web app named WebApp1. You use the WebJobs SDK to design a triggered App Service background task that automatically invokes a function in the code every time new data is received in a queue. You are preparing to configure the service processes a queue data item.

Which of the following is the service you should use?

- A. Logic Apps
- B. WebJobs**
- C. Flow
- D. Functions

Yes or No

QUESTION: 42

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing a medical records document management website. The website is used to store scanned copies of patient intake forms. If the stored intake forms are downloaded from storage by a third party, the contents of the forms must not be compromised. You need to store the intake forms according to the requirements.

Solution:

1. Create an Azure Cosmos DB database with Storage Service Encryption enabled.
2. Store the intake forms in the Azure Cosmos DB database.

Does the solution meet the goal?

- A. Yes
- B. No

QUESTION: 42

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing a medical records document management website. The website is used to store scanned copies of patient intake forms. If the stored intake forms are downloaded from storage by a third party, the contents of the forms must not be compromised. You need to store the intake forms according to the requirements.

Solution:

1. Create an Azure Cosmos DB database with Storage Service Encryption enabled.
2. Store the intake forms in the Azure Cosmos DB database.

Does the solution meet the goal?

- A. Yes
- B. No

Yes or No

QUESTION: 43

Your company has an Azure subscription. You need to deploy a number of Azure virtual machines to the subscription by using Azure Resource Manager (ARM) templates. The virtual machines will be included in a single availability set. You need to ensure that the ARM template allows for as many virtual machines as possible to remain accessible in the event of fabric failure or maintenance.

Which of the following is the value that you should configure for the platformFaultDomainCount property?

- A. 10
- B. 30
- C. Min Value
- D. Max Value

QUESTION: 43

Your company has an Azure subscription. You need to deploy a number of Azure virtual machines to the subscription by using Azure Resource Manager (ARM) templates. The virtual machines will be included in a single availability set. You need to ensure that the ARM template allows for as many virtual machines as possible to remain accessible in the event of fabric failure or maintenance.

Which of the following is the value that you should configure for the platformFaultDomainCount property?

- A. 10
- B. 30
- C. Min Value
- D. Max Value

QUESTION: 44

This question requires that you evaluate the underlined text to determine if it is correct.

You company has an on-premises deployment of MongoDB, and an Azure Cosmos DB account that makes use of the MongoDB API. You need to devise a strategy to migrate MongoDB to the Azure Cosmos DB account. You include the Data Management Gateway tool in your migration strategy.

Instructions: Review the underlined text. If it makes the statement correct, select 'No change required'. If the statement is incorrect, select the answer choice that makes the statement correct.

- A. No change required
- B. mongorestore
- C. Azure Storage Explorer
- D. AzCopy

QUESTION: 44

This question requires that you evaluate the underlined text to determine if it is correct.

You company has an on-premises deployment of MongoDB, and an Azure Cosmos DB account that makes use of the MongoDB API. You need to devise a strategy to migrate MongoDB to the Azure Cosmos DB account. You include the Data Management Gateway tool in your migration strategy.

Instructions: Review the underlined text. If it makes the statement correct, select 'No change required'. If the statement is incorrect, select the answer choice that makes the statement correct.

- A. No change required
- B. mongorestore
- C. Azure Storage Explorer
- D. AzCopy

Yes or No

QUESTION: 45

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing an application that processes Azure Blob storage events. Your application has the following requirements:

- Process transaction logs asynchronously for changes that occur to the blobs and the blob metadata.
- Process changes in the order in which they occurred.
- Retain changes for compliance reasons.

Solution: You use Azure Event Grid with a subscriber Azure Function app.

Does the solution meet the goal?

- A. Yes
- B. No

QUESTION: 45

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing an application that processes Azure Blob storage events. Your application has the following requirements:

- Process transaction logs asynchronously for changes that occur to the blobs and the blob metadata.
- Process changes in the order in which they occurred.
- Retain changes for compliance reasons.

Solution: You use Azure Event Grid with a subscriber Azure Function app.

Does the solution meet the goal?

- A. Yes
- B. No

QUESTION: 47

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing an application that processes Azure Blob storage events. Your application has the following requirements:

- Process transaction logs asynchronously for changes that occur to the blobs and the blob metadata.
- Process changes in the order in which they occurred.
- Retain changes for compliance reasons.

Solution: You use Azure Event Grid with a subscriber Azure Function app.

Does the solution meet the goal?

- A. Yes
- B. No

QUESTION: 47

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

You are developing an application that processes Azure Blob storage events. Your application has the following requirements:

- Process transaction logs asynchronously for changes that occur to the blobs and the blob metadata.
- Process changes in the order in which they occurred.
- Retain changes for compliance reasons.

Solution: You use Azure Event Grid with a subscriber Azure Function app.

Does the solution meet the goal?

- A. Yes
- B. No

Yes or No

QUESTION: 46

You are developing an Azure Cosmos DB solution by using the Azure Cosmos DB SQL API. The data includes millions of documents. Each document may contain hundreds of properties. The properties of the documents do not contain distinct values for partitioning. Azure Cosmos DB must scale individual containers in the database to meet the performance needs of the application by spreading the workload evenly across all partitions over time. You need to select a partition key.

Which two partition keys can you use? (Select two.)

- A. a single property value that does not appear frequently in the documents
- B. a value containing the collection name
- C. a single property value that appears frequently in the documents
- D. a concatenation of multiple property values with a random suffix appended
- E. a hash suffix appended to a property value

QUESTION: 46

You are developing an Azure Cosmos DB solution by using the Azure Cosmos DB SQL API. The data includes millions of documents. Each document may contain hundreds of properties. The properties of the documents do not contain distinct values for partitioning. Azure Cosmos DB must scale individual containers in the database to meet the performance needs of the application by spreading the workload evenly across all partitions over time. You need to select a partition key.

Which two partition keys can you use? (Select two.)

- A. a single property value that does not appear frequently in the documents
- B. a value containing the collection name
- C. a single property value that appears frequently in the documents
- D. a concatenation of multiple property values with a random suffix appended
- E. a hash suffix appended to a property value

QUESTION: 48

You are developing an e-Commerce Web App. You want to use Azure Key Vault to ensure that sign-ins to the e-Commerce Web App are secured by using Azure App Service authentication and Azure Active Directory (AAD).

What should you do on the e-Commerce Web App?

- A. Run the az keyvault secret command.
- B. Enable Azure AD Connect.
- C. Enable Managed Service Identity (MSI).
- D. Create an Azure AD service principal.

QUESTION: 48

You are developing an e-Commerce Web App. You want to use Azure Key Vault to ensure that sign-ins to the e-Commerce Web App are secured by using Azure App Service authentication and Azure Active Directory (AAD).

What should you do on the e-Commerce Web App?

- A. Run the az keyvault secret command.
- B. Enable Azure AD Connect.
- C. Enable Managed Service Identity (MSI).
- D. Create an Azure AD service principal.

QUESTION: 49

DRAG DROP

You have two Azure virtual networks named Hub1 and Spoke1. Hub1 connects to an on-premises network by using a Site-to-Site VPN connection. You are implementing peering between Hub1 and Spoke1. You need to ensure that a virtual machine connected to Spoke1 can connect to the on-premises network through Hub1.

How should you complete the PowerShell script? (To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Values

Answer Area

-AllowForwardedTraffic

```
$hub = Get-AZVirtualNetwork -ResourceGroup "RG1" -Name "Hub1"
```

-AllowGatewayTransit

```
$spoke = Get-AZVirtualNetwork -ResourceGroup "RG2" -Name "Spoke1"
```

-UseRemoteGateways

```
Add-AZVirtualNetworkPeering -Name "Hub1-Spoke1" -VirtualNetwork $hub
```

-RemoteVirtualNetworkId \$spoke.id

Value

```
Add-AZVirtualNetworkPeering -Name "Spoke1-Hub1" -VirtualNetwork $spoke
```

-RemoteVirtualNetworkId \$hub.id

Value

QUESTION: 49

DRAG DROP

You have two Azure virtual networks named Hub1 and Spoke1. Hub1 connects to an on-premises network by using a Site-to-Site VPN connection. You are implementing peering between Hub1 and Spoke1. You need to ensure that a virtual machine connected to Spoke1 can connect to the on-premises network through Hub1.

How should you complete the PowerShell script? (To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Values

Answer Area

-AllowForwardedTraffic

```
$hub = Get-AZVirtualNetwork -ResourceGroup "RG1" -Name "Hub1"
```

-AllowGatewayTransit

```
$spoke = Get-AZVirtualNetwork -ResourceGroup "RG2" -Name "Spoke1"
```

-UseRemoteGateways

```
-RemoteVirtualNetworkId $spoke.id      -AllowGatewayTransit
```

```
Add-AZVirtualNetworkPeering -Name "Hub1-Spoke1" -VirtualNetwork $hub
```

```
Add-AZVirtualNetworkPeering -Name "Spoke1-Hub1" -VirtualNetwork $spoke
```

-RemoteVirtualNetworkId \$hub.id

-UseRemoteGateways

QUESTION: 50

You are developing a web app that uses Azure Active Directory (Azure AD) for authentication. You want to configure the web app to use multifactor authentication.

What should you do?

- A. Enable mobile app authentication.
- B. In Azure AD conditional access, enable the baseline policy.
- C. In Azure AD, create a conditional access policy.
- D. Install the Azure Multi-Factor Authentication Server.

QUESTION: 50

You are developing a web app that uses Azure Active Directory (Azure AD) for authentication. You want to configure the web app to use multifactor authentication.

What should you do?

- A. Enable mobile app authentication.
- B. In Azure AD conditional access, enable the baseline policy.
- C. In Azure AD, create a conditional access policy.**
- D. Install the Azure Multi-Factor Authentication Server.

QUESTION: 51

You develop Azure solutions. You must connect to a No-SQL globally-distributed database by using the .NET API. You need to create an object to configure and execute requests in the database.

Which code segment should you use?

- A. new Container(EndpointUri, PrimaryKey);
- B. new Database(EndpointUri, PrimaryKey);
- C. new CosmosClient(EndpointUri, PrimaryKey);

QUESTION: 51

You develop Azure solutions. You must connect to a No-SQL globally-distributed database by using the .NET API. You need to create an object to configure and execute requests in the database.

Which code segment should you use?

- A. new Container(EndpointUri, PrimaryKey);
- B. new Database(EndpointUri, PrimaryKey);
- C. new CosmosClient(EndpointUri, PrimaryKey);

QUESTION: 52

This question requires that you evaluate the underlined text to determine if it is correct.

Your Azure Active Directory Azure (Azure AD) tenant has an Azure subscription linked to it. Your developer has created a mobile application that obtains Azure AD access tokens using the OAuth 2 implicit grant type. The mobile application must be registered in Azure AD. You require a redirect URL from the developer for registration purposes.

Instructions: Review the underlined text. If it makes the statement correct, select 'No change is needed. If the statement is incorrect, select the answer choice that makes the statement correct.

- A. No change required.
- B. a secret
- C. a login hint
- D. a client ID

QUESTION: 52

This question requires that you evaluate the underlined text to determine if it is correct.

Your Azure Active Directory Azure (Azure AD) tenant has an Azure subscription linked to it. Your developer has created a mobile application that obtains Azure AD access tokens using the OAuth 2 implicit grant type. The mobile application must be registered in Azure AD. You require a redirect URL from the developer for registration purposes.

Instructions: Review the underlined text. If it makes the statement correct, select 'No change is needed. If the statement is incorrect, select the answer choice that makes the statement correct.

- A. No change required.
- B. a secret
- C. a login hint
- D. a client ID

QUESTION: 53

You have an Azure virtual network that contains a subnet named Subnet1. Subnet1 is associated to a network security group (NSG) named NSG1. NSG1 blocks all outbound traffic that is not allowed explicitly. Subnet1 contains virtual machines that must communicate with the Azure Cosmos DB service. You need to create an outbound security rule in NSG1 to enable the virtual machines to connect to Azure Cosmos DB.

What should you include in the solution?

- A. a service tag
- B. a private endpoint
- C. a subnet delegation
- D. an application security group

QUESTION: 53

You have an Azure virtual network that contains a subnet named Subnet1. Subnet1 is associated to a network security group (NSG) named NSG1. NSG1 blocks all outbound traffic that is not allowed explicitly. Subnet1 contains virtual machines that must communicate with the Azure Cosmos DB service. You need to create an outbound security rule in NSG1 to enable the virtual machines to connect to Azure Cosmos DB.

What should you include in the solution?

- A. a service tag
- B. a private endpoint
- C. a subnet delegation
- D. an application security group

QUESTION: 54

You are creating an Azure key vault using PowerShell. Objects deleted from the key vault must be kept for a set period of 90 days.

Which two of the following parameters must be used in conjunction to meet the requirement?
(Select two.)

- A. EnabledForDeployment
- B. EnablePurgeProtection
- C. EnabledForTemplateDeployment
- D. EnableSoftDelete

QUESTION: 54

You are creating an Azure key vault using PowerShell. Objects deleted from the key vault must be kept for a set period of 90 days.

Which two of the following parameters must be used in conjunction to meet the requirement?
(Select two.)

- A. EnabledForDeployment
- B. EnablePurgeProtection
- C. EnabledForTemplateDeployment
- D. EnableSoftDelete

Yes or No

QUESTION: 55

Note: This question is part of a series of questions that present the same scenario. Each question in the series contains a unique solution that might meet the stated goals. Some question sets might have more than one correct solution, while others might not have a correct solution.

Your company has an Azure Active Directory (Azure AD) environment. Users occasionally connect to Azure AD via the Internet. You need to ensure that users who connect to Azure AD via the internet using an unidentified IP address, are automatically instructed to change their passwords.

Solution: You configure the use of Azure Key Vault.

Does the solution meet the goal?

A. Yes

B. No

Yes or No

QUESTION: 56

You are developing a Java application that uses Cassandra to store key and value data. You plan to use a new Azure Cosmos DB resource and the Cassandra API in the application. You create an Azure Active Directory (Azure AD) group named Cosmos DB Creators to enable provisioning of Azure Cosmos accounts, databases, and containers. The Azure AD group must not be able to access the keys that are required to access the data. You need to restrict access to the Azure AD group.

Which role-based access control should you use?

- A. DocumentDB Accounts Contributor
- B. Cosmos Backup Operator
- C. Cosmos DB Operator
- D. Cosmos DB Account Reader

QUESTION: 56

You are developing a Java application that uses Cassandra to store key and value data. You plan to use a new Azure Cosmos DB resource and the Cassandra API in the application. You create an Azure Active Directory (Azure AD) group named Cosmos DB Creators to enable provisioning of Azure Cosmos accounts, databases, and containers. The Azure AD group must not be able to access the keys that are required to access the data. You need to restrict access to the Azure AD group.

Which role-based access control should you use?

- A. DocumentDB Accounts Contributor
- B. Cosmos Backup Operator
- C. Cosmos DB Operator
- D. Cosmos DB Account Reader

QUESTION: 57

You develop a Web App on a tier D1 app service plan. You notice that page load times increase during periods of peak traffic. You want to implement automatic scaling when CPU load is above 80 percent. Your solution must minimize costs.

What should you do first?

- A. Enable autoscaling on the Web App.
- B. Switch to the Premium App Service tier plan.
- C. Switch to the Standard App Service tier plan.
- D. Switch to the Azure App Services consumption plan.

QUESTION: 57

You develop a Web App on a tier D1 app service plan. You notice that page load times increase during periods of peak traffic. You want to implement automatic scaling when CPU load is above 80 percent. Your solution must minimize costs.

What should you do first?

- A. Enable autoscaling on the Web App.
- B. Switch to the Premium App Service tier plan.
- C. Switch to the Standard App Service tier plan.
- D. Switch to the Azure App Services consumption plan.

QUESTION: 58

Your company's Azure subscription includes an Azure Log Analytics workspace. Your company has a hundred on-premises servers that run either Windows Server 2012 R2 or Windows Server 2016, and is linked to the Azure Log Analytics workspace. The Azure Log Analytics workspace is set up to gather performance counters associated with security from these linked servers. You must configure alerts based on the information gathered by the Azure Log Analytics workspace. You have to make sure that alert rules allow for dimensions, and that alert creation time should be kept to a minimum. Furthermore, a single alert notification must be created when the alert is created and when the alert is resolved. You need to make use of the necessary signal type when creating the alert rules.

Which of the following is the option you should use?

- A. The Activity log signal type.
- B. The Application Log signal type.
- C. The Metric signal type.
- D. The Audit Log signal type.

QUESTION: 58

Your company's Azure subscription includes an Azure Log Analytics workspace. Your company has a hundred on-premises servers that run either Windows Server 2012 R2 or Windows Server 2016, and is linked to the Azure Log Analytics workspace. The Azure Log Analytics workspace is set up to gather performance counters associated with security from these linked servers. You must configure alerts based on the information gathered by the Azure Log Analytics workspace. You have to make sure that alert rules allow for dimensions, and that alert creation time should be kept to a minimum. Furthermore, a single alert notification must be created when the alert is created and when the alert is resolved. You need to make use of the necessary signal type when creating the alert rules.

Which of the following is the option you should use?

- A. The Activity log signal type.
- B. The Application Log signal type.
- C. The Metric signal type.
- D. The Audit Log signal type.

QUESTION: 59

You accept an IT internship at a local charity. The charity wants you to help them with compliance and auditing requirements. You need to ensure that a column or combination of columns uniquely identifies each row of a table.

Which constraint should you define?

- A. Primary key
- B. Secondary key
- C. Foreign key
- D. Default key

QUESTION: 59

You accept an IT internship at a local charity. The charity wants you to help them with compliance and auditing requirements. You need to ensure that a column or combination of columns uniquely identifies each row of a table.

Which constraint should you define?

- A. Primary key
- B. Secondary key
- C. Foreign key
- D. Default key

QUESTION: 60

The terms "bitmap," "b-tree," and "hash" refer to which type of database structure?

- A. View
- B. Function
- C. Index
- D. Stored procedure
- E. Trigger

QUESTION: 60

The terms "bitmap," "b-tree," and "hash" refer to which type of database structure?

- A. View
- B. Function
- C. Index
- D. Stored procedure
- E. Trigger

QUESTION: 61

This question requires that you evaluate the underlined text to determine if it is correct.

The CREATE TABLE command removes one or more table definitions and all data, indexes, triggers, constraints, and permission specifications for those tables.

Review the text above. If it makes the statement correct, select "No change is needed." If the statement is incorrect, select the answer choice that makes the statement correct

- A. No change is needed
- B. DROP TABLE
- C. TRUNCATE TABLE
- D. ALTER TABLE

QUESTION: 61

This question requires that you evaluate the underlined text to determine if it is correct.

The CREATE TABLE command removes one or more table definitions and all data, indexes, triggers, constraints, and permission specifications for those tables.

Review the text above. If it makes the statement correct, select "No change is needed." If the statement is incorrect, select the answer choice that makes the statement correct

- A. No change is needed
- B. DROP TABLE**
- C. TRUNCATE TABLE
- D. ALTER TABLE

QUESTION: 62

You are a developer at your company. You need to update the definitions for an existing Logic App.

What should you use?

- A. the Enterprise Integration Pack (EIP)
- B. the Logic App Code View
- C. the API Connections
- D. the Logic Apps Designer

QUESTION: 62

You are a developer at your company. You need to update the definitions for an existing Logic App.

What should you use?

- A. the Enterprise Integration Pack (EIP)
- B. the Logic App Code View**
- C. the API Connections
- D. the Logic Apps Designer

QUESTION: 63

You are developing an Azure App Service REST API. The API must be called by an Azure App Service web app. The API must retrieve and update user profile information stored in Azure Active Directory (Azure AD). You need to configure the API to make the updates.

Which two tools should you use? (Select two.)

- A. Microsoft Graph API
- B. Microsoft Authentication Library (MSAL)
- C. Azure API Management
- D. Microsoft Azure Security Center
- E. Microsoft Azure Key Vault SDK

QUESTION: 64

You have an Azure virtual network and an on-premises datacenter. You are planning a Site-to-Site VPN connection between the datacenter and the virtual network.

Which two resources should you include in your plan? (Select two.)

- A. a user-defined route
- B. a virtual network gateway
- C. Azure Firewall
- D. Azure Web Application Firewall (WAF)
- E. an on-premises data gateway
- F. an Azure application gateway
- G. a local network gateway

QUESTION: 64

You have an Azure virtual network and an on-premises datacenter. You are planning a Site-to-Site VPN connection between the datacenter and the virtual network.

Which two resources should you include in your plan? (Select two.)

- A. a user-defined route
- B. a virtual network gateway
- C. Azure Firewall
- D. Azure Web Application Firewall (WAF)
- E. an on-premises data gateway
- F. an Azure application gateway
- G. a local network gateway

QUESTION: 65

Your company has an on-premises network and three Azure subscriptions named Subscription1, Subscription2, and Subscription3. The departments at the company use the Azure subscriptions as shown in the following table.

Department	Subscription
IT	Subscription1
Research	Subscription1
Development	Subscription2
Testing	Subscription2
Distribution	Subscription3

All the resources in the subscriptions are in either the West US Azure region or the West US 2 Azure region. You plan to connect all the subscriptions to the on-premises network by using ExpressRoute.

What is the minimum number of ExpressRoute circuits required?

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5

What is the minimum number of ExpressRoute circuits required?

- A. 1
- B. 2
- C. 3
- D. 4
- E. 5

QUESTION: 66

You are developing a .NET Core MVC application that allows customers to research independent holiday accommodation providers. You want to implement Azure Search to allow the application to search the index by using various criteria to locate documents related to accommodation. You want the application to allow customers to search the index by using regular expressions.

What should you do?

- A. Configure the SearchMode property of the SearchParameters class.
- B. Configure the QueryType property of the SearchParameters class.
- C. Configure the Facets property of the SearchParameters class.
- D. Configure the Filter property of the SearchParameters class.

QUESTION: 66

You are developing a .NET Core MVC application that allows customers to research independent holiday accommodation providers. You want to implement Azure Search to allow the application to search the index by using various criteria to locate documents related to accommodation. You want the application to allow customers to search the index by using regular expressions.

What should you do?

- A. Configure the SearchMode property of the SearchParameters class.
- B. Configure the QueryType property of the SearchParameters class.**
- C. Configure the Facets property of the SearchParameters class.
- D. Configure the Filter property of the SearchParameters class.

QUESTION: 67

You have an Azure subscription that contains 100 virtual machines. You plan to design a data protection strategy to encrypt the virtual disks. You need to recommend a solution to encrypt the disks by using Azure Disk Encryption. The solution must provide the ability to encrypt operating system disks and data disks.

What should you include in the recommendation?

- A. a certificate
- B. a key
- C. a passphrase
- D. a secret

QUESTION: 67

You have an Azure subscription that contains 100 virtual machines. You plan to design a data protection strategy to encrypt the virtual disks. You need to recommend a solution to encrypt the disks by using Azure Disk Encryption. The solution must provide the ability to encrypt operating system disks and data disks.

What should you include in the recommendation?

- A. a certificate
- B. a key**
- C. a passphrase
- D. a secret

QUESTION: 68

You have two Azure Data Factory instances named ADFdev and ADFprod. ADFdev connects to an Azure DevOps Git repository. You publish changes from the main branch of the Git repository to ADFdev. You need to deploy the artifacts from ADFdev to ADFprod.

What should you do first?

- A. From ADFdev, modify the Git configuration.
- B. From ADFdev, create a linked service.
- C. From Azure DevOps, create a release pipeline.
- D. From Azure DevOps, update the main branch.

QUESTION: 68

You have two Azure Data Factory instances named ADFdev and ADFprod. ADFdev connects to an Azure DevOps Git repository. You publish changes from the main branch of the Git repository to ADFdev. You need to deploy the artifacts from ADFdev to ADFprod.

What should you do first?

- A. From ADFdev, modify the Git configuration.
- B. From ADFdev, create a linked service.
- C. From Azure DevOps, create a release pipeline.**
- D. From Azure DevOps, update the main branch.

QUESTION: 69

HOTSPOT

You are designing a real-time dashboard solution that will visualize streaming data from remote sensors that connect to the internet. The streaming data must be aggregated to show the average value of each 10-second interval. The data will be discarded after being displayed in the dashboard. The solution will use Azure Stream Analytics and must meet the following requirements:

- Minimize latency from an Azure Event hub to the dashboard.
- Minimize the required storage.
- Minimize development effort.

Answer Area

Azure Stream Analytics input type:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

Azure Stream Analytics output type:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

Aggregation query location:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

Answer Area

Azure Stream Analytics input type:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

Azure Stream Analytics output type:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

Aggregation query location:

Azure Event Hub
Azure SQL Database
Azure Stream Analytics
Microsoft Power BI

QUESTION: 70

You are developing a solution that will stream to Azure Stream Analytics. The solution will have both streaming data and reference data.

Which input type should you use for the reference data?

- A. Azure Cosmos DB
- B. Azure Blob storage
- C. Azure IoT Hub
- D. Azure Event Hubs

QUESTION: 70

You are developing a solution that will stream to Azure Stream Analytics. The solution will have both streaming data and reference data.

Which input type should you use for the reference data?

- A. Azure Cosmos DB
- B. Azure Blob storage**
- C. Azure IoT Hub
- D. Azure Event Hubs

QUESTION: 71

You are designing an Azure Stream Analytics job to process incoming events from sensors in retail environments. You need to process the events to produce a running average of shopper counts during the previous 15 minutes, calculated at five-minute intervals.

Which type of window should you use?

- A. snapshot
- B. tumbling
- C. hopping
- D. sliding

QUESTION: 71

You are designing an Azure Stream Analytics job to process incoming events from sensors in retail environments. You need to process the events to produce a running average of shopper counts during the previous 15 minutes, calculated at five-minute intervals.

Which type of window should you use?

- A. snapshot
- B. tumbling
- C. hopping
- D. sliding

QUESTION: 72

You have data files in Azure Blob storage. You plan to transform the files and move them to Azure Data Lake Storage. You need to transform the data by using mapping data flow.

Which Azure service should you use?

- A. Azure Data Box Gateway
- B. Azure Storage Sync
- C. Azure Data Factory
- D. Azure Databricks

QUESTION: 72

You have data files in Azure Blob storage. You plan to transform the files and move them to Azure Data Lake Storage. You need to transform the data by using mapping data flow.

Which Azure service should you use?

- A. Azure Data Box Gateway
- B. Azure Storage Sync
- C. Azure Data Factory
- D. Azure Databricks

QUESTION: 73

Your company has offices in New York and Amsterdam. The company has an Azure subscription. Both offices connect to Azure by using a Site-to-Site VPN connection. The office in Amsterdam uses resources in the North Europe Azure region. The office in New York uses resources in the East US Azure region. You need to implement ExpressRoute circuits to connect each office to the nearest Azure region. Once the ExpressRoute circuits are connected, the on-premises computers in the Amsterdam office must be able to connect to the on-premises servers in the New York office by using the ExpressRoute circuits.

Which ExpressRoute option should you use?

- A. ExpressRoute FastPath
- B. ExpressRoute Global Reach
- C. ExpressRoute Direct
- D. ExpressRoute Local

QUESTION: 73

Your company has offices in New York and Amsterdam. The company has an Azure subscription. Both offices connect to Azure by using a Site-to-Site VPN connection. The office in Amsterdam uses resources in the North Europe Azure region. The office in New York uses resources in the East US Azure region. You need to implement ExpressRoute circuits to connect each office to the nearest Azure region. Once the ExpressRoute circuits are connected, the on-premises computers in the Amsterdam office must be able to connect to the on-premises servers in the New York office by using the ExpressRoute circuits.

Which ExpressRoute option should you use?

- A. ExpressRoute FastPath
- B. ExpressRoute Global Reach**
- C. ExpressRoute Direct
- D. ExpressRoute Local

QUESTION: 74

You have an Azure virtual machine named VM1 that runs Windows Server 2019 and contains 500 GB of data files. You are designing a solution that will use Azure Data Factory to transform the data files, and then load the files to Azure Data Lake Storage.

What should you deploy on VM1 to support the design?

- A. the Azure Pipelines agent
- B. the Azure File Sync agent
- C. the On-premises data gateway
- D. the self-hosted integration runtime in Azure

QUESTION: 74

You have an Azure virtual machine named VM1 that runs Windows Server 2019 and contains 500 GB of data files. You are designing a solution that will use Azure Data Factory to transform the data files, and then load the files to Azure Data Lake Storage.

What should you deploy on VM1 to support the design?

- A. the Azure Pipelines agent
- B. the Azure File Sync agent
- C. the On-premises data gateway
- D. the self-hosted integration runtime in Azure

QUESTION: 75

You have an app named App1 that uses two on-premises Microsoft SQL Server databases named DB1 and DB2. You plan to migrate DB1 and DB2 to Azure. You need to recommend an Azure solution to host DB1 and DB2.

The solution must meet the following requirements:

- Support server-side transactions across DB1 and DB2.
- Minimize administrative effort to update the solution.

What should you recommend?

- A. two Azure SQL databases in an elastic pool
- B. two Azure SQL databases on different Azure SQL Database servers
- C. two Azure SQL databases on the same Azure SQL Database managed instance
- D. two SQL Server databases on an Azure virtual machine

QUESTION: 75

You have an app named App1 that uses two on-premises Microsoft SQL Server databases named DB1 and DB2. You plan to migrate DB1 and DB2 to Azure. You need to recommend an Azure solution to host DB1 and DB2.

The solution must meet the following requirements:

- Support server-side transactions across DB1 and DB2.
- Minimize administrative effort to update the solution.

What should you recommend?

- A. two Azure SQL databases in an elastic pool
- B. two Azure SQL databases on different Azure SQL Database servers
- C. two Azure SQL databases on the same Azure SQL Database managed instance
- D. two SQL Server databases on an Azure virtual machine

QUESTION: 76

While attending college, you accept an IT internship at a local charity. The charity needs to report on data that is related and exists in two tables. You need to establish a relationship between the data that is in the two tables.

Which constraint should you define?

- A. Foreign key
- B. Index key
- C. Link key
- D. Default key

QUESTION: 76

While attending college, you accept an IT internship at a local charity. The charity needs to report on data that is related and exists in two tables. You need to establish a relationship between the data that is in the two tables.

Which constraint should you define?

- A. Foreign key
- B. Index key
- C. Link key
- D. Default key

QUESTION: 77

You have a table named Product that contains one million rows. You need to search for product information in the Product table by using the product's unique ID.

What will make this type of search more efficient?

- A. A cursor
- B. A subquery
- C. A trigger
- D. An index

QUESTION: 77

You have a table named Product that contains one million rows. You need to search for product information in the Product table by using the product's unique ID.

What will make this type of search more efficient?

- A. A cursor
- B. A subquery
- C. A trigger
- D. An index

QUESTION: 78

You have the following table definition:

```
CREATE TABLE Product  
(ProductID INTEGER,  
Name VARCHAR(20))
```

You need to insert a new product. The product's name is Plate and the product's ID is 12345.

Which statement should you use?

- A.

```
INSERT INTO Product
VALUES (ProductID = 12345, Name = 'Plate')
```

- B.

```
INSERT NEW ProductID = 12345, Name = 'Plate'
INTO Product
```

- C.

```
INSERT 12345, 'Plate'
INTO Product
```

- D.

```
INSERT INTO Product (ProductID, Name)
VALUES (12345, 'Plate')
```

- A. Option A
- B. Option B
- C. Option C
- D. Option D

- A.

```
INSERT INTO Product
VALUES (ProductID = 12345, Name = 'Plate')
```

 - B.

```
INSERT NEW ProductID = 12345, Name = 'Plate'
INTO Product
```

 - C.

```
INSERT 12345, 'Plate'
INTO Product
```

 - D.

```
INSERT INTO Product (ProductID, Name)
VALUES (12345, 'Plate')
```
-
- A. Option A
 - B. Option B
 - C. Option C
 - D. Option D

QUESTION: 79

You have a table that contains information about all students in your school.

Which SQL keyword should you use to change a student's first name in the table?

- A. UPDATE
- B. CHANGE
- C. SELECT
- D. INSERT

QUESTION: 79

You have a table that contains information about all students in your school.

Which SQL keyword should you use to change a student's first name in the table?

- A. UPDATE
- B. CHANGE
- C. SELECT
- D. INSERT

QUESTION: 80

You are designing a microservices architecture that will support a web application. The solution must meet the following requirements:

- Allow independent upgrades to each microservice.
- Deploy the solution on-premises and to Azure.
- Set policies for performing automatic repairs to the microservices.
- Support low-latency and hyper-scale operations.

You need to recommend a technology.

- A. Azure Container Instance
- B. Azure Virtual Machine Scale Set
- C. Azure Service Fabric
- D. Azure Logic App

QUESTION: 80

You are designing a microservices architecture that will support a web application. The solution must meet the following requirements:

- Allow independent upgrades to each microservice.
- Deploy the solution on-premises and to Azure.
- Set policies for performing automatic repairs to the microservices.
- Support low-latency and hyper-scale operations.

You need to recommend a technology.

- A. Azure Container Instance
- B. Azure Virtual Machine Scale Set
- C. Azure Service Fabric
- D. Azure Logic App

QUESTION: 81

You assign User1 a set of permissions that include the WITH GRANT OPTION.

The WITH GRANT OPTION enables User1 to:

- A. request a log of permission use.
- B. delegate permissions to other users.
- C. create new database users.
- D. view other users' permissions.

QUESTION: 81

You assign User1 a set of permissions that include the WITH GRANT OPTION.

The WITH GRANT OPTION enables User1 to:

- A. request a log of permission use.
- B. delegate permissions to other users.**
- C. create new database users.
- D. view other users' permissions.

QUESTION: 82

Which type of index changes the order in which the data is stored in a table?

- A. non-sequential
- B. sequential
- C. non-clustered
- D. clustered

QUESTION: 82

Which type of index changes the order in which the data is stored in a table?

- A. non-sequential
- B. sequential
- C. non-clustered
- D. clustered

QUESTION: 83

You use Azure virtual machines to run a custom application that uses an Azure SQL Database instance on the back end. The IT department at your company recently enabled forced tunneling. Since the configuration change, developers have noticed degraded performance when they access the database from the Azure virtual machine. You need to recommend a solution to minimize latency when accessing the database. The solution must minimize costs.

What should you include in the recommendation?

- A. Virtual Network (VNET) service endpoints
- B. Azure virtual machines that run Microsoft SQL Server servers
- C. Azure SQL Database Managed Instance
- D. Always On availability groups

QUESTION: 83

You use Azure virtual machines to run a custom application that uses an Azure SQL Database instance on the back end. The IT department at your company recently enabled forced tunneling. Since the configuration change, developers have noticed degraded performance when they access the database from the Azure virtual machine. You need to recommend a solution to minimize latency when accessing the database. The solution must minimize costs.

What should you include in the recommendation?

- A. Virtual Network (VNET) service endpoints
- B. Azure virtual machines that run Microsoft SQL Server servers
- C. Azure SQL Database Managed Instance
- D. Always On availability groups

QUESTION: 84

A company has a SaaS solution that uses Azure SQL Database with elastic pools. The solution contains a dedicated database for each customer organization. Customer organizations have peak usage at different periods during the year. You need to implement the Azure SQL Database elastic pool to minimize cost.

Which option or options should you configure?

- A. Number of transactions only
- B. eDTUs per database only
- C. Number of databases only
- D. CPU usage only
- E. eDTUs and max data size

QUESTION: 84

A company has a SaaS solution that uses Azure SQL Database with elastic pools. The solution contains a dedicated database for each customer organization. Customer organizations have peak usage at different periods during the year. You need to implement the Azure SQL Database elastic pool to minimize cost.

Which option or options should you configure?

- A. Number of transactions only
- B. eDTUs per database only
- C. Number of databases only
- D. CPU usage only
- E. eDTUs and max data size

QUESTION: 85

A company manages several on-premises Microsoft SQL Server databases. You need to migrate the databases to Microsoft Azure by using a backup and restore process.

Which data technology should you use?

- A. Azure SQL Database single database
- B. Azure SQL Data Warehouse
- C. Azure Cosmos DB
- D. Azure SQL Database Managed Instance

QUESTION: 85

A company manages several on-premises Microsoft SQL Server databases. You need to migrate the databases to Microsoft Azure by using a backup and restore process.

Which data technology should you use?

- A. Azure SQL Database single database
- B. Azure SQL Data Warehouse
- C. Azure Cosmos DB
- D. Azure SQL Database Managed Instance

QUESTION: 86

You have an Azure subscription. Users access the resources in the subscription from either home or from customer sites. From home, users must establish a point-to-site VPN to access the Azure resources. The users on the customer sites access the Azure resources by using site-to-site VPNs. You have a line-of-business app named App1 that runs on several Azure virtual machine. The virtual machines run Windows Server 2016. You need to ensure that the connections to App1 are spread across all the virtual machines.

What are possible Azure services that you can use?

- A. a public load balancer, an Azure Application Gateway
- B. Traffic Manager, an Azure Application Gateway
- C. an Azure Content Delivery Network (CDN)
- D. an internal load balancer, an Azure Application Gateway

QUESTION: 86

You have an Azure subscription. Users access the resources in the subscription from either home or from customer sites. From home, users must establish a point-to-site VPN to access the Azure resources. The users on the customer sites access the Azure resources by using site-to-site VPNs. You have a line-of-business app named App1 that runs on several Azure virtual machine. The virtual machines run Windows Server 2016. You need to ensure that the connections to App1 are spread across all the virtual machines.

What are possible Azure services that you can use?

- A. a public load balancer, an Azure Application Gateway
- B. Traffic Manager, an Azure Application Gateway
- C. an Azure Content Delivery Network (CDN)
- D. an internal load balancer, an Azure Application Gateway

QUESTION: 87

You plan to create a speech recognition deep learning model. The model must support the latest version of Python. You need to recommend a deep learning framework for speech recognition to include in the Data Science Virtual Machine (DSVM).

What should you recommend?

- A. Rattle
- B. TensorFlow
- C. Weka
- D. Deeplearning4j

QUESTION: 87

You plan to create a speech recognition deep learning model. The model must support the latest version of Python. You need to recommend a deep learning framework for speech recognition to include in the Data Science Virtual Machine (DSVM).

What should you recommend?

- A. Rattle
- B. TensorFlow**
- C. Weka
- D. Deeplearning4j

QUESTION: 88

You plan to use a Deep Learning Virtual Machine (DLVM) to train deep learning models using Compute Unified Device Architecture (CUDA) computations. You need to configure the DLVM to support CUDA.

What should you implement?

- A. Solid State Drives (SSD)
- B. Computer Processing Unit (CPU) speed increase by using overclocking
- C. Graphic Processing Unit (GPU)
- D. High Random Access Memory (RAM) configuration
- E. Intel Software Extensions (Intel SGX) technology

QUESTION: 88

You plan to use a Deep Learning Virtual Machine (DLVM) to train deep learning models using Compute Unified Device Architecture (CUDA) computations. You need to configure the DLVM to support CUDA.

What should you implement?

- A. Solid State Drives (SSD)
- B. Computer Processing Unit (CPU) speed increase by using overclocking
- C. Graphic Processing Unit (GPU)**
- D. High Random Access Memory (RAM) configuration
- E. Intel Software Extensions (Intel SGX) technology

QUESTION: 89

Which counter should you monitor for real-time processing to meet the technical requirements?

- A. SU% Utilization
- B. CPU% utilization
- C. Concurrent users
- D. Data Conversion Errors

QUESTION: 89

Which counter should you monitor for real-time processing to meet the technical requirements?

- A. SU% Utilization
- B. CPU% utilization**
- C. Concurrent users
- D. Data Conversion Errors

QUESTION: 90
HOTSPOT

You are evaluating the role assignments.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
DBAGroup1 will be able to sign in to each customer's Azure SQL database by using Azure Data Studio.	<input type="radio"/>	<input type="radio"/>
DBAGroup1 will be able to assign the SQL DB Contributor role to other users.	<input type="radio"/>	<input type="radio"/>
DBAGroup2 will be able to create a new Azure SQL database on each customer's Azure SQL Database server.	<input type="radio"/>	<input type="radio"/>

QUESTION: 90
HOTSPOT

You are evaluating the role assignments.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Answer Area

Statements	Yes	No
DBAGroup1 will be able to sign in to each customer's Azure SQL database by using Azure Data Studio.	<input checked="" type="radio"/>	<input type="radio"/>
DBAGroup1 will be able to assign the SQL DB Contributor role to other users.	<input type="radio"/>	<input checked="" type="radio"/>
DBAGroup2 will be able to create a new Azure SQL database on each customer's Azure SQL Database server.	<input checked="" type="radio"/>	<input type="radio"/>

QUESTION: 91

You have an Azure SQL managed instance. You need to gather the last execution of a query plan and its runtime statistics. The solution must minimize the impact on currently running queries.

What should you do?

- A. Generate an estimated execution plan.
- B. Generate an actual execution plan.
- C. Run sys.dm_exec_query_plan_stats.
- D. Generate Live Query Statistics.

QUESTION: 91

You have an Azure SQL managed instance. You need to gather the last execution of a query plan and its runtime statistics. The solution must minimize the impact on currently running queries.

What should you do?

- A. Generate an estimated execution plan.
- B. Generate an actual execution plan.
- C. Run `sys.dm_exec_query_plan_stats`.
- D. Generate Live Query Statistics.

QUESTION: 92

Denormalization is performed in order to:

- A. Reduce redundancy.
- B. Eliminate repeating groups.
- C. Create smaller tables.
- D. Improve query performance.

QUESTION: 92

Denormalization is performed in order to:

- A. Reduce redundancy.
- B. Eliminate repeating groups.
- C. Create smaller tables.
- D. Improve query performance.

QUESTION: 93

In SQL, an insert statement is used to add a:

- A. User to a database.
- B. Row of data to a table.
- C. Table to a database.
- D. Column to a table definition.

QUESTION: 93

In SQL, an insert statement is used to add a:

- A. User to a database.
- B. Row of data to a table.**
- C. Table to a database.
- D. Column to a table definition.

QUESTION: 94

Your database contains a table named Customer. You need to delete the record from the Customer table that has a CustomerID of 12345.

Which statement should you use?

- A. UPDATE Customer
 DELETE *
 WHERE CustomerID = 12345
- B. DELETE CustomerID
 FROM Customer
 WHERE CustomerID = 12345
- C. UPDATE CustomerID
 FROM Customer
 DELETE *
 WHERE CustomerID = 12345
- D. DELETE FROM Customer
 WHERE CustomerID = 12345

- A. Option A
- B. Option B
- C. Option C
- D. Option D

QUESTION: 94

Your database contains a table named Customer. You need to delete the record from the Customer table that has a CustomerID of 12345.

Which statement should you use?

- A. UPDATE Customer
 DELETE *
 WHERE CustomerID = 12345
- B. DELETE CustomerID
 FROM Customer
 WHERE CustomerID = 12345
- C. UPDATE CustomerID
 FROM Customer
 DELETE *
 WHERE CustomerID = 12345
- D. DELETE FROM Customer
 WHERE CustomerID = 12345

- A. Option A
- B. Option B
- C. Option C
- D. Option D

QUESTION: 95

You have a SQL Server database named MyDB that uses SQL Server Authentication.

Which connection string should you use to connect to MyDB?

- A. Data Source=MyDB; UserID=username; Password=P@sswOrd; Initial Catalog=Sales
- B. Data Source=MyDB; Integrated Security=SSPI; Initial Catalog=Sales
- C. Data Source=MyDB; Integrated Security=True; Initial Catalog=Sales
- D. Data Source=MyDB; Trusted_Connection=True; MultipleActiveResultSets=True; Initial Catalog=Sales

QUESTION: 95

You have a SQL Server database named MyDB that uses SQL Server Authentication.

Which connection string should you use to connect to MyDB?

- A. Data Source=MyDB; UserID=username; Password=P@sswOrd; Initial Catalog=Sales
- B. Data Source=MyDB; Integrated Security=SSPI; Initial Catalog=Sales
- C. Data Source=MyDB; Integrated Security=True; Initial Catalog=Sales
- D. Data Source=MyDB; Trusted_Connection=True; MultipleActiveResultSets=True; Initial Catalog=Sales

QUESTION: 96

You are developing a database that other programmers will query to display race results. You need to provide the ability to query race results without allowing access to other information in the database.

What should you do?

- A. Disable implicit transactions.
- B. Place the query into a stored procedure.
- C. Create an index on the result table.
- D. Add an AFTER UPDATE trigger on the result table to reject updates.

QUESTION: 96

You are developing a database that other programmers will query to display race results. You need to provide the ability to query race results without allowing access to other information in the database.

What should you do?

- A. Disable implicit transactions.
- B. Place the query into a stored procedure.**
- C. Create an index on the result table.
- D. Add an AFTER UPDATE trigger on the result table to reject updates.